

# Time Series Forecasting with Neural Networks

M.Sc. Economics  
University of Bologna  
A.Y. 2020/2021

- Accurate time series forecasting is crucial in both economic and financial applications
- Standard econometric approaches rely almost exclusively on *linear models* (AR, MA, ARIMA...) as they usually perform well on standard tasks.
- Developments in deep learning have made available highly flexible models that can infer any complex nonlinearity from the data (**Universal Approximation Theorem**)

- Existing literature is lacking a unified perspective on the performance of neural networks on economic and financial forecasting tasks.
- Only a few types of neural networks (typically two) are usually tested against each other. Here, we compare *five different neural network models*.
- The difference between *univariate* and *multivariate* performance is not always assessed
- Economic and financial series are never contrasted with more 'standard' (well-behaved) series

The **linear models** estimated in this project are:

- The **Random Walk (RW)**, a simple *persistence model*:

$$y_t = y_{t-1} + \epsilon_t$$

where  $\epsilon_t$  is a *white noise process* (zero mean and constant variance).

- The **Autoregressive Integrated Moving Average (ARIMA)**, taking the form:

$$\Delta y_t = \alpha + \phi_1 \Delta y_{t-1} + \dots + \phi_p \Delta y_{t-p} + \theta_1 \epsilon_{t-1} + \dots + \theta_q \epsilon_{t-q} + \epsilon_t$$

where:

- **p** is the number of *autoregressive components*
- **d** is the *order of integration* ( $\Delta$ )
- **q** is the number of *moving average components*

# Feedforward Neural Networks - 1/2

A **feedforward neural network** (FFNN) is a sequence of **dense layers** taking the form:


$$f(X) = \beta_0 + \sum_{k=1}^K \beta_k h_k(X)$$

where  $K$  is the number of **neurons**, each triggered by an **activation**  $h_k$ :

$$h_k(X) = g \left( w_{k0} + \sum_{j=1}^p w_{kj} X_j \right)$$

where  $g(\cdot)$  is a **nonlinear activation function** (e.g. ReLU, tanh, sigmoid) and  $p$  are the predictors.

# Feedforward Neural Networks - 2/2



**Figure:** Single-layer Feedforward Network with two neurons  
(Extended (left) and schematic (right) representations)

Convolutional Neural Networks (CNNs) make use of *two specific layers*:

- **Convolution Layer**
- **Max Pooling Layer**

# Convolutional Neural Networks - 2/4

The **convolution layer** uses a **kernel (k)** and *convolves* it over the **input data (X)**, producing a **feature map (s)**:

$$X = [0.4, 0.1, 0.5, 0.3, 0.9] \quad k = [1, 0, 1]$$



$$s = [0.4 + 0 + 0.5, 0.1 + 0 + 0.3, 0.5 + 0 + 0.9] = [0.9, 0.4, 1.4]$$



The **max pooling layer** performs *dimension reduction* selecting the highest value within a given **pool size**:

$$s = [0.9, 0.4, 1.4] \quad \text{pool size} = 2$$



$$s = [\max(0.9, 0.4), \max(0.4, 1.4)] = [0.9, 1.4]$$

# Convolutional Neural Networks - 4/4




Figure: Convolutional Neural Network for 1-d Time Series Forecasting

- Recurrent neural networks (RNNs) are designed to model *sequences of data*.
- They take in a sequence  $X = \{X_1, X_2, \dots, X_T\}$  and pass it through a sequence of **hidden states**  $\{h_t\}_1^T = \{h_1, h_2, \dots, h_T\}$ :

$$h_{tk} = g \left( w_{k0} + \sum_{j=1}^p u_{kj} X_{tj} + \sum_{s=1}^K w_{ks} h_{t-1,s} \right)$$

for  $K$  neurons in each hidden state.

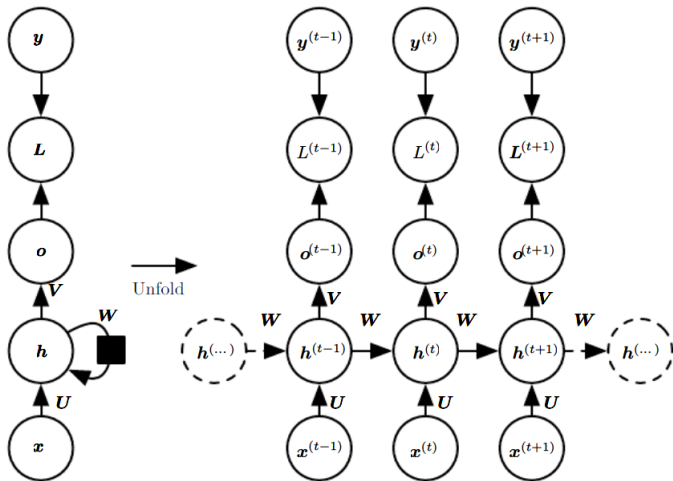
- Note that the value of the present state is **dependent on the value of the previous state**.

- The **output** of the RNN at each time step is a weighted sum of the hidden states' neurons

$$O_t = \beta_0 + \sum_{k=1}^K \beta_k h_{tk}$$

- The weights in each layer are **not time-dependent**

# Recurrent Neural Networks - 3/5




**Figure:** Single-layer Recurrent Network  
(Schematic (left) and extended (right) representations)

# Recurrent Neural Networks - 4/5

The **Long Short-Term Memory (LSTM)** network is a type of RNN where each cell is made up of:


- **Forget Gate** → regulates how much past information is discarded and how much is kept
- **Input Gate** → updates the cell state
- **Output Gate** → defines the value of the next hidden state



# Recurrent Neural Networks - 5/5

The **Gated Recurrent Unit (GRU)** network is a type of RNN where each cell consists of:

- **Update Gate** → decides how much new information to add
- **Reset Gate** → decides how much old information to discard

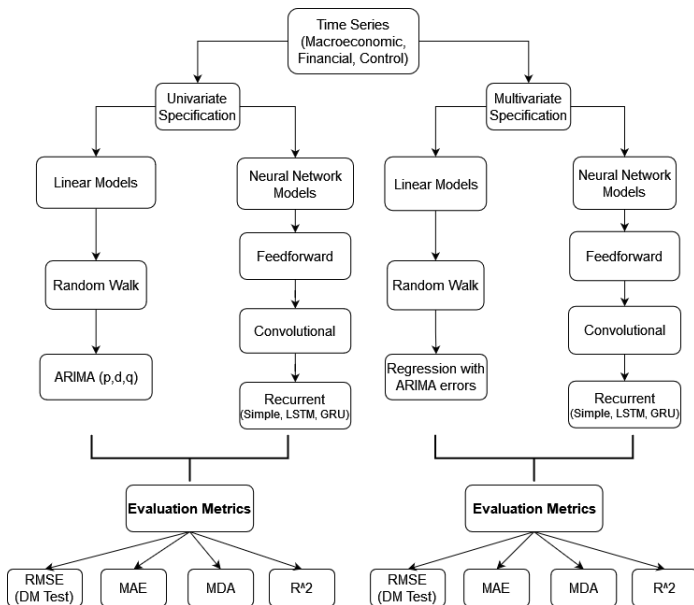


**Three time series** were used in this forecasting exercise:

- **Macroeconomic Time Series**: *YoY % change in unemployment in the United States between July 1955 and December 2019 (monthly observations)*  
**Covariates**: *Inflation rate, YoY % change in Federal funds rate and industrial production*
- **Financial Time Series**: *S&P 500 Trading Volume between January 4th, 2010 to December 30th, 2020 (daily observations)*  
**Covariates**: *Daily FTSE 100 and Nasdaq trading volumes, CBOE Volatility Index*
- **Control Time Series**: *Daily visits to a statistical forecasting website from September 14, 2014 to August 19, 2020*  
**Covariates**: *Unique Visits, First Time Visits and Returning Visits to the same website*



# Methodology - 2/4





- Each series was split into *training* and *test* sets.
- Fitting of the ARIMA models was carried out using the **auto\_arma** function from Python's *pmdarima*
- Neural Networks' **hyperparameters** were selected using a **grid search algorithm** and the models were optimized using **TensorFlow**

- Each model produces **one-step ahead forecasts** on both the training (in-sample) and test (out-of-sample) sets
- **Walk-forward validation** was used with linear models, while neural networks performed **sliding-window regression**
- To avoid overfitting, the number of hidden layers is **built conditionally on training set size** (larger dataset = deeper network)



## Univariate

- The ARIMA model has statistically significant lower RMSE **in-sample**
- There is no statistically significant difference between ARIMA and neural networks RMSE **out-of-sample**
- CNN yields a **31% increase in MDA with respect to the ARIMA model** and an **average 34% increase in MDA with respect to other neural networks** (MDA = 0.59)





## Multivariate

- Neural network models have statistically significant lower RMSE **in-sample**
- Again, there is no statistically significant difference between ARIMA and neural networks RMSE **out-of-sample**
- CNN outperforms all other models on MDA





## Univariate

- ARIMA has a significantly lower in-sample RMSE
- Neural networks perform **at least as good as ARIMA** out-of-sample
- GRU yields a **low but statistically significant reduction in RMSE with respect to ARIMA**
- CNN yields a **17% increase in MDA with respect to the ARIMA model**





## Multivariate

- Neural network models perform **significantly worse** than ARIMAX on all available metrics



## Univariate



- **Neural networks outperform ARIMA** in every metric
- **Lowest** reduction in RMSE is **53%**, **highest** is **55%** (FFNN)
- **Lowest** increase in MDA is **30%** (CNN), **highest** is **34%** (LSTM)





## Multivariate

- ARIMAX performs significantly better than neural networks on **all in-sample and out-of-sample metrics** ( $R^2 = .99$ )
- Neural network out-of-sample performance is *slightly worse* than in the univariate specification



Possible improvements on the present approach include:

- Check performance on **different forecasting horizons**
- Operate some **transformation of the input data**
- Widen the grid-search or change **hyperparameter tuning algorithm**
- Check performance of **hybrid models** (e.g. ARIMA-ANN)

- When properly trained, neural networks are able to reach the accuracy of linear models (and sometimes even outperform them) in the **univariate setting**.
- Evidence in the **multivariate setting** is more in favour of linear models.
- CNNs achieve **high accuracy on the MDA measure** in both macroeconomic and financial series.
- **RNNs are not as accurate as expected** in time series forecasting tasks.