Time Series Forecasting with Neural Networks

Defended by: Supervisor:

M.Sc. Economics University of Bologna A.Y. 2020/2021

Motivation - 1/2

- Accurate time series forecasting is crucial in both economic and financial applications
- Standard econometric approaches rely almost exclusively on linear models (AR, MA, ARIMA...) as they usually perform well on standard tasks.
- Developments in deep learning have made available highly flexible models that can infer any complex nonlinearity from the data (Universal Approximation Theorem)

Motivation - 2/2

- Existing literature is lacking a unified perspective on the performance of neural networks on economic and financial forecasting tasks.
- Only a few types of neural networks (typically two) are usually tested against each other. Here, we compare five different neural network models.
- The difference between *univariate* and *multivariate* performance is not always assessed
- Economic and financial series are never contrasted with more 'standard' (well-behaved) series

Linear Models

The **linear models** estimated in this project are:

• The Random Walk (RW), a simple persistence model:

$$y_t = y_{t-1} + \epsilon_t$$

where ϵ_t is a white noise process (zero mean and constant variance).

 The Autoregressive Integrated Moving Average (ARIMA), taking the form:

$$\Delta y_t = \alpha + \phi_1 \Delta y_{t-1} + \ldots + \phi_p \Delta y_{t-p} + \theta_1 \epsilon_{t-1} + \ldots + \theta_q \epsilon_{t-q} + \epsilon_t$$

where:

- p is the number of autoregressive components
- **d** is the order of integration (Δ)
- **q** is the number of moving average components

Feedforward Neural Networks - 1/2

A **feedforward neural network** (FFNN) is a sequence of **dense layers** taking the form:

$$f(X) = \beta_0 + \sum_{k=1}^K \beta_k h_k(X)$$

where K is the number of **neurons**, each triggered by an **activation** h_k :

$$h_k(X) = g\left(w_{k0} + \sum_{j=1}^p w_{kj}X_j\right)$$

where $g(\cdot)$ is a **nonlinear activation function** (e.g. ReLU, tanh, sigmoid) and p are the predictors.

Feedforward Neural Networks - 2/2

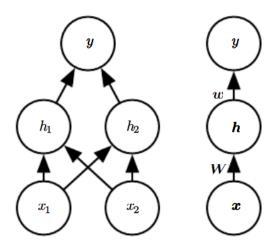


Figure: Single-layer Feedforward Network with two neurons (Extended (left) and schematic (right) representations)

Convolutional Neural Networks - 1/4

Convolutional Neural Networks (CNNs) make use of two specific layers:

- Convolution Layer
- Max Pooling Layer

Convolutional Neural Networks - 2/4

The **convolution layer** uses a **kernel** (**k**) and *convolves* it over the **input data** (**X**), producing a **feature map** (**s**):

$$X = [0.4, 0.1, 0.5, 0.3, 0.9]$$
 $k = [1, 0, 1]$

$$\downarrow \downarrow s = [0.4 + 0 + 0.5, 0.1 + 0 + 0.3, 0.5 + 0 + 0.9] = [0.9, 0.4, 1.4]$$

Convolutional Neural Networks - 3/4

The **max pooling layer** performs *dimension reduction* selecting the highest value within a given **pool size**:

$$s = [0.9, 0.4, 1.4]$$
 pool size = 2
 \Downarrow
 $s = [\max(0.9, 0.4), \max(0.4, 1.4)] = [0.9, 1.4]$

Convolutional Neural Networks - 4/4

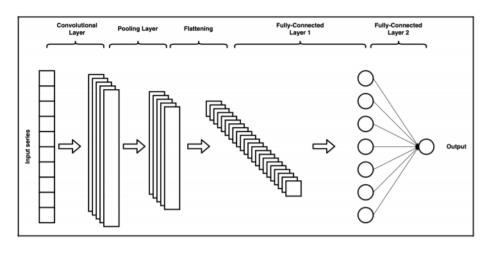


Figure: Convolutional Neural Network for 1-d Time Series Forecasting

Recurrent Neural Networks - 1/5

- Recurrent neural networks (RNNs) are designed to model sequences of data.
- They take in a sequence $X = \{X_1, X_2, \dots, X_T\}$ and pass it through a sequence of **hidden states** $\{h_t\}_1^T = \{h_1, h_2, \dots, h_T\}$:

$$h_{tk} = g \left(w_{k0} + \sum_{j=1}^{p} u_{kj} X_{tj} + \sum_{s=1}^{K} w_{ks} h_{t-1,s} \right)$$

for K neurons in each hidden state.

 Note that the value of the present state is dependent on the value of the previous state.

Recurrent Neural Networks - 2/5

 The output of the RNN at each time step is a weighted sum of the hidden states' neurons

$$O_t = \beta_0 + \sum_{k=1}^K \beta_k h_{tk}$$

• The weights in each layer are not time-dependent

Recurrent Neural Networks - 3/5

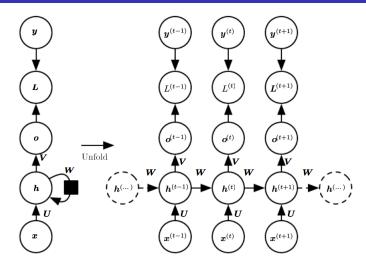
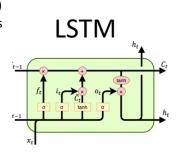


Figure: Single-layer Recurrent Network (Schematic (left) and extended (right) representations)

Recurrent Neural Networks - 4/5

The Long Short-Term Memory (LSTM) network is a type of RNN where each cell is made up of:

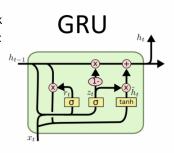
- Forget Gate → regulates how much past information is discarded and how much is kept
- ullet Input Gate ightarrow updates the cell state
- Output Gate → defines the value of the next hidden state



Recurrent Neural Networks - 5/5

The **Gated Recurrent Unit (GRU)** network is a type of RNN where each cell consists of:

- Update Gate → decides how much new information to add
- Reset Gate → decides how much old information to discard



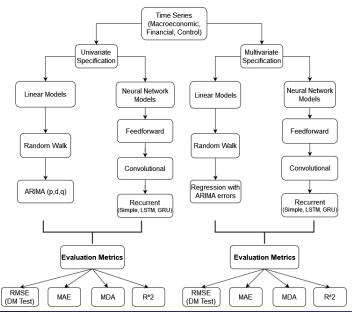
Methodology - 1/4

Three time series were used in this forecasting exercise:

- Macroeconomic Time Series: YoY % change in unemployment in the United States between July 1955 and December 2019 (monthly observations)
 - **Covariates**: Inflation rate, YoY % change in Federal funds rate and industrial production
- Financial Time Series: S&P 500 Trading Volume between January 4th, 2010 to December 30th, 2020 (daily observations)

 Covariates: Daily FTSE 100 and Nasdaq trading volumes, CBOE Volatility Index
- Control Time Series : Daily visits to a statistical forecasting website from September 14, 2014 to August 19, 2020
 Covariates: Unique Visits, First Time Visits and Returning Visits to the same website

Methodology - 2/4



Methodology - 3/4

- Each series was split into training and test sets.
- Fitting of the ARIMA models was carried out using the **auto_arima** function from Python's *pmdarima*
- Neural Networks' hyperparameters were selected using a grid search algorithm and the models were optimized using TensorFlow

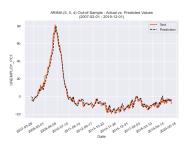
Methodology - 4/4

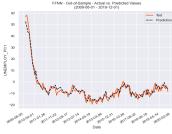
- Each model produces one-step ahead forecasts on both the training (in-sample) and test (out-of-sample) sets
- Walk-forward validation was used with linear models, while neural networks performed sliding-window regression
- To avoid overfitting, the number of hidden layers is built conditionally on training set size (larger dataset = deeper network)

Results - Macroeconomic Time Series - 1/2

Univariate

- The ARIMA model has statistically significant lower RMSE in-sample
- There is no statistically significant difference between ARIMA and neural networks RMSE out-of-sample
- CNN yields a 31% increase in MDA with respect to the ARIMA model and an average 34% increase in MDA with respect to other neural networks (MDA = 0.59)

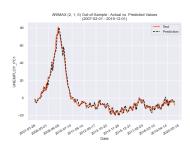


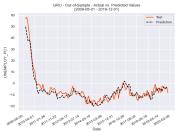


Results - Macroeconomic Time Series - 2/2

Multivariate

- Neural network models have statistically significant lower RMSE in-sample
- Again, there is no statistically significant difference between ARIMA and neural networks RMSE out-of-sample
- CNN outperforms all other models on MDA

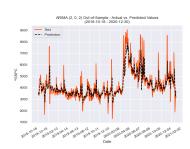


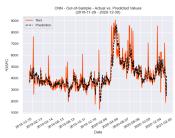


Results - Financial Time Series - 1/2

Univariate

- ARIMA has a significantly lower in-sample RMSE
- Neural networks perform at least as good as ARIMA out-of-sample
- GRU yields a low but statistically significant reduction in RMSE with respect to ARIMA
- CNN yields a 17% increase in MDA with respect to the ARIMA model

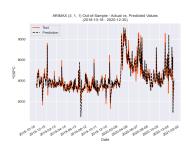


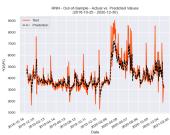


Results - Financial Time Series - 2/2

Multivariate

 Neural network models perform significantly worse than ARIMAX on all available metrics

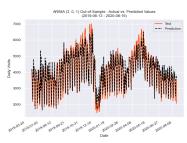


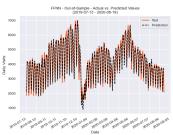


Results - Control Time Series - 1/2

Univariate

- Neural networks outperform ARIMA in every metric
- Lowest reduction in RMSE is 53%, highest is 55% (FFNN)
- Lowest increase in MDA is 30% (CNN), highest is 34% (LSTM)

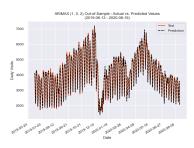


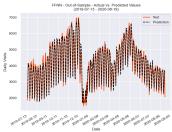


Results - Control Time Series - 2/2

Multivariate

- ARIMAX performs significantly better than neural networks on all in-sample and out-of-sample metrics (R² = .99)
- Neural network out-of-sample performance is slightly worse than in the univariate specification





Future Work

Possible improvements on the present approach include:

- Check performance on different forecasting horizons
- Operate some transformation of the input data
- Widen the grid-search or change hyperparameter tuning algorithm
- Check performance of hybrid models (e.g. ARIMA-ANN)

Conclusions

- When properly trained, neural networks are able to reach the accuracy of linear models (and sometimes even outperform them) in the univariate setting.
- Evidence in the multivariate setting is more in favour of linear models.
- CNNs achieve high accuracy on the MDA measure in both macroeconomic and financial series.
- RNNs are not as accurate as expected in time series forecasting tasks.